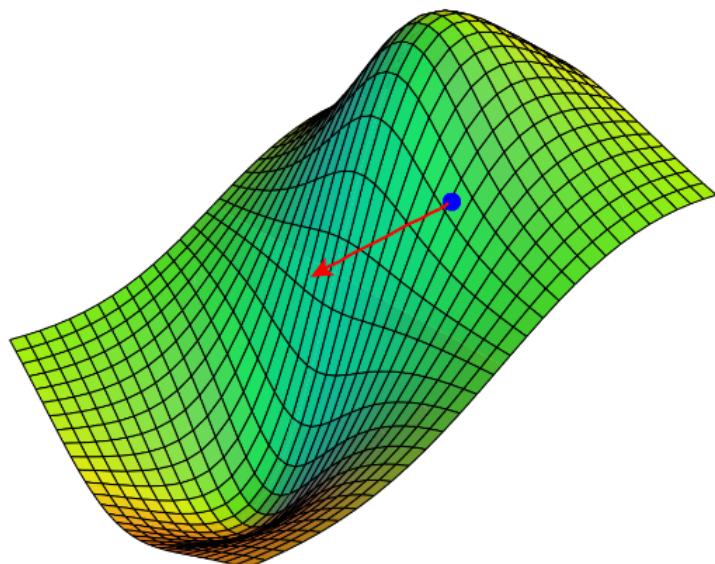


Quantitative Research on Market Predictability: Machine Learning and the Efficient Market Hypothesis



Aurouet Lucas

Master EKAP
IAE Nantes
Under the supervision of Prof. O.DARNE 01/03/2021

Contents

1	Introduction	2
2	Theory	3
2.1	Statistics & Machine Learning	3
2.1.1	The Perceptron: a simple example to understand how machines learn	3
2.1.2	Networks & Backpropagation	7
2.1.3	Activation functions and enhanced gradient descent	11
2.1.4	Support Vector Machine & the kernel trick	11
2.1.5	Soft Margin SVM	14
2.2	Finance	14
3	Application	14
3.1	Data analysis	14
4	Conclusion	26

1 Introduction

Trading is futile. This is an oversimplification of the Efficient Market Hypothesis. According to Eugene Fama who formulated this hypothesis in 1970¹, prices at a certain date incorporate all the available and relevant information at that date. There are never overvalued nor undervalued assets, fundamental and/or technical analysis are therefore pointless as all the signals are already reflected by prices on the market. One cannot beat the market, is another way of expressing the EMH, which brings us to our initial statement; trading is futile. As trading basically consists in predicting an asset performance, under the EMH, trading is a wasted effort. It is quite hard to understand how, if the EMH holds, financial institutions have poured significant amounts of money into trying to outperform the market. Even less understandable is how some have succeeded. I think it is reasonable to define what 'outperforming the market' means. 'Outperforming the market' encompasses the achievement of greater returns than a portfolio containing every asset available on the market, with the same level of risk offered by such portfolio. Alternatively, it could mean achieving the same return, with less risk. Eugene Fama states that the only way to expect returns greater than the market is to take additional risks, or that you cannot achieve that same return on investment with a risk inferior to the market risk. We are now able to ask the following question: does the EMH holds with real life data ? Is trading futile ? Can you achieve greater returns than the market for a corresponding market risk ? Or is one constantly better off holding a large index like the S&P 500 for instance ? In my opinion, questioning the EMH is the most riveting topic in quantitative finance as it answers questions from a very theoretical point of view, typically what is observed in academics but also from a very practical point of view, the kind we observe on a trading floor. Econometrics is the science of combining these two aspects, using theoretical tools to shine lights on very down to earth problems. Therefore this work will try to reflect both dimensions. There are many ways to prove or disprove the EMH, but consistently beating the market is one of them (consistency is essential here). The objective here is to use data science to try and outperform the market. The results should definitely not be taken as an attempt to generate profits but rather as an exercise where the objective is to make use of quantitative techniques to elucidate a question. Machine Learning is a very powerful tool when a problem comes to identifying patterns and regularities in the data. As we will discuss later on, it is particularly adapted in our case. One fact often played down is that Machine Learning boils down to a list of algorithms applied in a particular order, and it is very easy to make mistakes when using automated techniques. We hope to have sufficient knowledge to question our own results, especially when working in quantitative finance. Although we are not, as mentioned above, attempting to concretely apply our findings, We think it is only natural to add some perspective when working with financial data. Finance is connected to the economy by many pipelines and misleading Artificial Intelligence can induce disastrous consequences. Artificial Intelligence is a very broad term and Machine Learning is one of its component, the later relies on the idea of being able to find patterns without explicitly defining these patterns, whereas Artificial Intelligence relies on the idea of writing a program capable of the same thought process as a human. They can sometimes be interchangeable and the frontier between the two is often blurry, here we are focusing on learning patterns from data, hence we will almost exclusively refer to Machine Learning as opposed to Artificial intelligence. We will also refer to Statistical Learning, which could be seen as the gap between traditional statistical

¹Efficient Capital Markets: A Review of Theory and Empirical Work - E.Fama

models and Machine Learning models. Statistical Learning solves data related problems by formulating them in a formal statistical problem which will then be solved using parts of Machine Learning algorithms. This gives us a large panel of tools to extract useful information from the data, coerce this information into a functional form that can then be exploited to infer the behavior of unseen data. We will first gaze over the Statistical and Machine Learning methods we will be using ranging from penalized regression to neural networks, we will then detail the financial side of the problem, going over technical analysis, fundamental analysis and the necessary financial mathematics. The second part is a step-by-step journal of the progress of the overall project and what modifications have been made to improve the data processing².

This project is coded in Python 3.8 and is associated with a Jupyter Notebook, available on [GitHub](#). The integrity of the code is open source.

I often compare open source to science. To where science took this whole notion of developing ideas in the open and improving on other peoples' ideas and making it into what science is today and the incredible advances that we have had. - Linus Torvalds, developer of Linux and Git

2 Theory

2.1 Statistics & Machine Learning

2.1.1 The Perceptron: a simple example to understand how machines learn

Artificial Neural Networks (ANN or NN for short) are Machine Learning models designed to replicate the way the human brain processes information and is heavily inspired by biology. In a human brain, the information travels through a network of neurons which can activate to pass the signal to the next neurons. Neurons are connected by branches in which the information transits, starting from a particular point, these branches, along with the neurons that activated create a path to an end point which outputs a decision. The choice for each neuron to carry on the information is what determines the final output. ANN replicate this logic by creating a stream of branches and neurons inside which the information is processed to eventually reach a particular output.

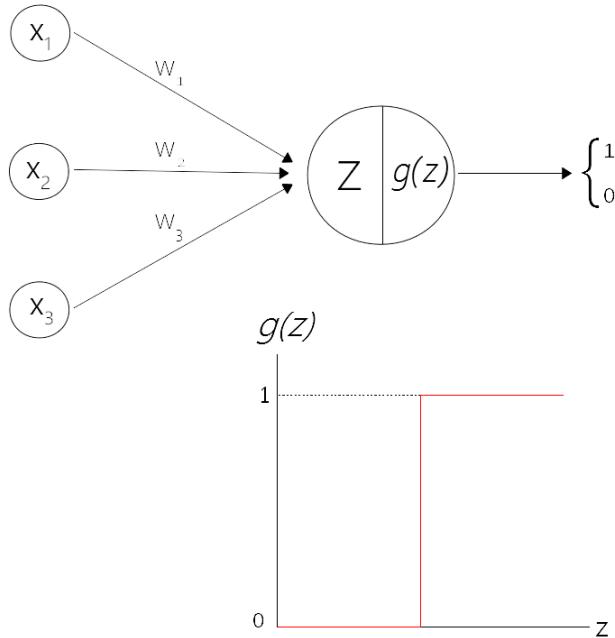
In this section, we will focus on the Perceptron, a particular ANN which consists of only one neuron. It is easier to understand the logic behind ANN using a Perceptron, the results can then be generalized to more complex networks with more neurons. We will try to understand how a Perceptron is able to find patterns in the data without being explicitly programmed. This is what we refer to as the learning algorithm, or how the machine learns from the data.

We begin by giving the model inputs (x_p ³) that account for the exogenous variables. Each input is multiplied by a weight w_j and all weighted inputs are summed. The

²Although data processing often refers to the pre-screening of the data before feeding it into a model, here we refer to the entire process from the data gathering to the diagnosis of the results

³p = 1,..., P where P is the total number of inputs

Figure 1: Diagram of a Perceptron



weighted sum, z goes through an activation function $g(z)$, that outputs 1 if the sum exceeds a certain threshold, 0 otherwise. For instance, we can think of inputs as values for technical indicators such as the RSI, the MACD and the stochastic oscillator and the output as 1 if we should "buy" and 0 if we should "sell". We take any data point and extract the values for the inputs and for the output, we compute the weighted sum and if that number exceeds the threshold, the Perceptron will output 1, or "buy".

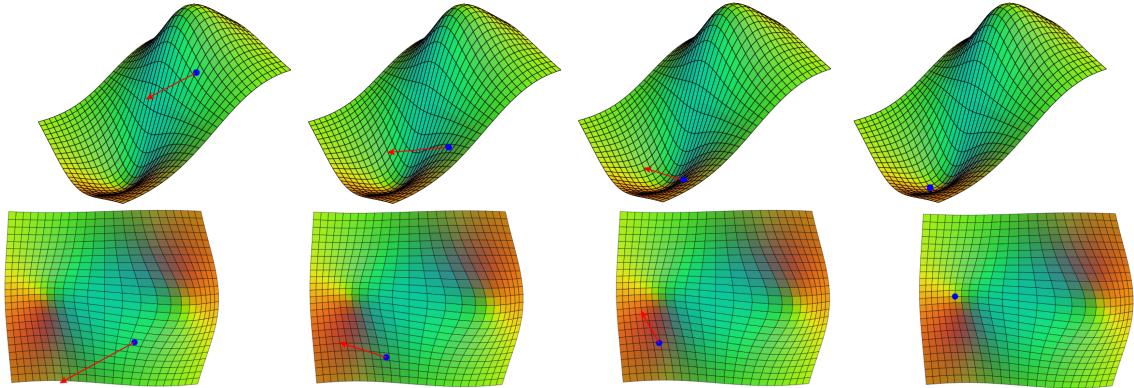
To learn how to accurately classify data, the Perceptron will update each weight such that we eventually reach a point where every data point is associated to the correct class. If the Perceptron outputs "buy", and the correct position was indeed buying, we can repeat the procedure for another data point. If the correct position was to sell, we update the weights and then we move on to the next data point. Weights are updated whenever the Perceptron outputs the wrong value, and kept the same whenever the Perceptron outputs the correct value. As weights are the only elements that can be manipulated, by constantly updating them we will eventually come to the right combination of weights, which outputs the correct buy or sell output for every data point. The weights are then no longer updated and the Perceptron has finished training. For each new set of values for the technical indicators, the weights will be such that the weighted sum of inputs activate the neuron only when necessary, or when buying is the right call, and stay deactivated when the right position is to sell. The weights update is not random, we desire to increment weights in order to increase the accuracy at each update. In order to do so, we can calculate a loss function, the loss will be 0 if the output corresponds to the target value and different from 0 if the Perceptron fails. This loss function represents the error, or how accurate the Perceptron is, with respect to every weights combination. We can minimize this loss function by differentiating it with respect to each weights, which solves for the weights that outputs the minimal error⁴. This closely relates to Ordinary Least Squares where we minimize the Mean Squared Error in order to find coefficients that minimize the

⁴0 ideally, but in reality the minimal error lies somewhere between 0 and 1

squared error. Although with OLS, one can find the solution to the differential equation analytically, it is often impossible with Perceptrons, and ANN, as they do not have closed-form solutions. To find the correct weights we are required to use numerical methods, and update them gradually towards the minimum instead of directly deriving weights that minimize the loss function.

Gradient descent is a numerical algorithm to help find the local minimum of a function (denoted as C), assuming it is differentiable. Starting with random initial values for the weights we compute the vector of partial derivatives of the loss function with respect to every weights, also called the gradient vector of the loss function ($\nabla(C)$). This vector of partial derivatives will point in the direction of steepest ascent where any increment in the weights results in the largest possible step on the curve. If for a particular set of weights we land at a point a on the loss function, the gradient tells us in which direction the loss function increases the most or which new combination of weights will land on a' as far up as possible from a . As our objective is to find the minimum of the loss function, we take the inverse of the gradient vector, which then leads to the the direction of steepest descent. The inverse gradient tells us how to update weights so that the loss function decreases the most rapidly towards a local minimum. Once we updated the weights, we repeat the procedure and compute the gradient vector of the loss function evaluated at a the new weights combination. We take a step in the direction of steepest descent (inverse of the gradient vector at that point) and this procedure is iterated until we reach a minimum. At a local minimum all the partial derivatives are 0, hence the gradient vector is a null vector. Less formally, once we reach a local minimum, the gradient descent tells us that we should move by 0 in every direction, which is equivalent to not moving at all.

Figure 2: Gradient Descent



For each data point we compute the error and the gradient at that point, take the inverse and increment the weights in the direction pointed by the inverse gradient. By always choosing the direction of steepest descent we will eventually converge to a local minimum, given that the step size is adapted.

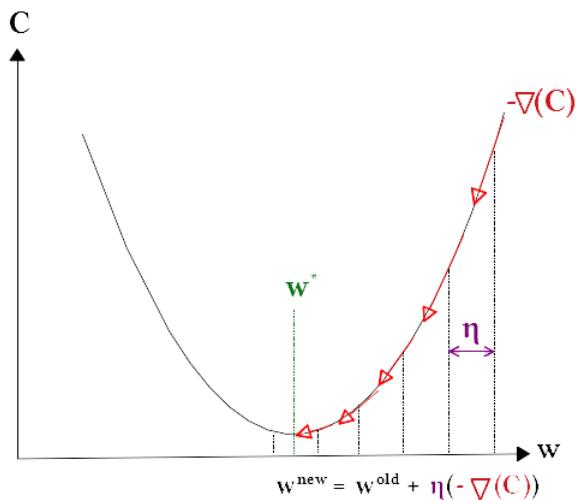
$$\overrightarrow{w^{new}} = \overrightarrow{w^{old}} + \eta(-\nabla(C)) \quad (1)$$

Where $\overrightarrow{w^{new}}$ is the vector of updated weights, $\overrightarrow{w^{old}}$ is the vector of old weights,

η is the learning rate which determines the size of the step and $\nabla(C) = \begin{bmatrix} \frac{\partial C}{\partial w_j} \\ \dots \\ \frac{\partial C}{\partial w_J} \end{bmatrix}$ is the vector of partial derivatives of the loss function (C) with respect to each weights in $w^{\overrightarrow{old}}$.

We often determine stopping rules that stop the gradient descent algorithm instead of reaching a point where all partial derivatives are 0. The algorithm is efficient nonetheless, the gradient vector does not point towards the minimum but the direction in which the loss function decreases. If we take a step in that direction once we are really close from the minimum, we might overshoot and land beyond the minimum, therefore if no stopping rules are applied the algorithm keeps iterating indefinitely. Once we reach a certain number of iterations, or that we are satisfied with the model's accuracy, we can voluntarily stop the procedure and get values for the weights that minimize the loss function. It is a possibility that the minimum identified by the gradient descent algorithm is a local minimum, while we ideally want to find the global minimum of the loss function. By construction, gradient descent is not able to tell whether we attained a local or global minimum which can lead to significant inaccuracies in the model. This is one of the crucial drawbacks in gradient descent.

Figure 3: Overshooting the local minimum

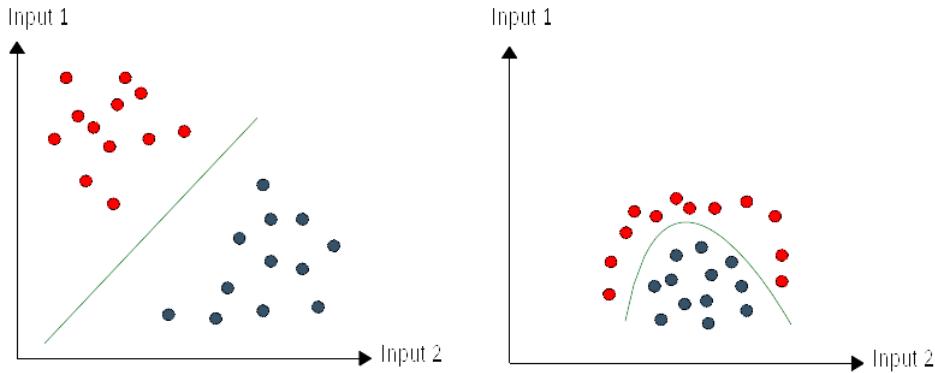


In the simple example above, we see how inappropriate step size η can lead the algorithm to oscillate around the local minimum of the loss function which is found at weight w^* . The algorithm will never converge as weight will "bounce" from right to left, when the weight is to large the gradient points towards the left and because the step size is to large, we land beyond the minimum. The weight is then to small, the gradient points towards the right and we take a step in that direction landing beyond w^* . The algorithm will iterate indefinitely as we never reach w^* where $\nabla(C)$ is null. This shows the importance of determining stopping rules such as number of iterations or minimum acceptable accuracy to prevent infinite iterations.

The Perceptron is limited to linearly separable problems, where it is possible to find a linear decision frontier in the input space that separates classes. In two dimensions, the frontier is a straight line, in three dimensions it would be a plane and for all further dimensions the frontier is an hyperplane. Once the neuron has been trained and we

the final weights are obtained we can calculate $b(x) = W^T x + w_0$ to trace a decision boundary. Every point falling on a particular side of the boundary will be associated with the corresponding class and any point falling on the other side of the frontier will be associated to the other class. The diagram below illustrates two different problems, the first one which is linearly separable and the second one which is not. Because Perceptron is a linear classifier, we need to upgrade it in order to solve non linear problems such as the one presented on the diagram above.

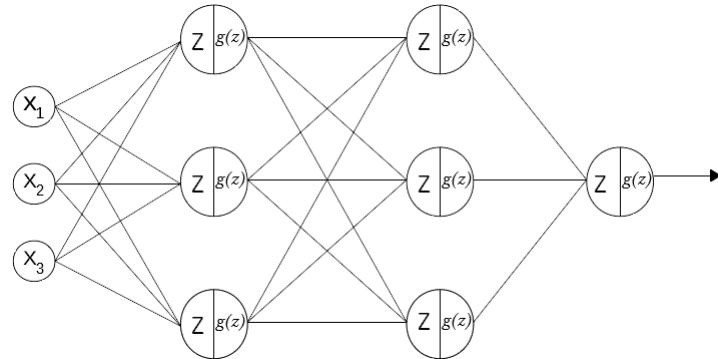
Figure 4: Linear separability



2.1.2 Networks & Backpropagation

The use of Neural Networks allows us to find non-linear decision boundaries and separate classes that could not be identified by a Perceptron. Neural Networks are a series of Perceptrons interconnected with each other. They consist in one entry layer, one output layer and intermediate "hidden layers" whose size may vary. The size of a layer is determined by the number of neurons and the ultimate size of the network is determined by the size of each layer multiplied by the number of layers. The following example is a Neural Network with one entry layer of 3 neurons, two hidden layers of 3 neurons each and one output layer of 1 neuron.

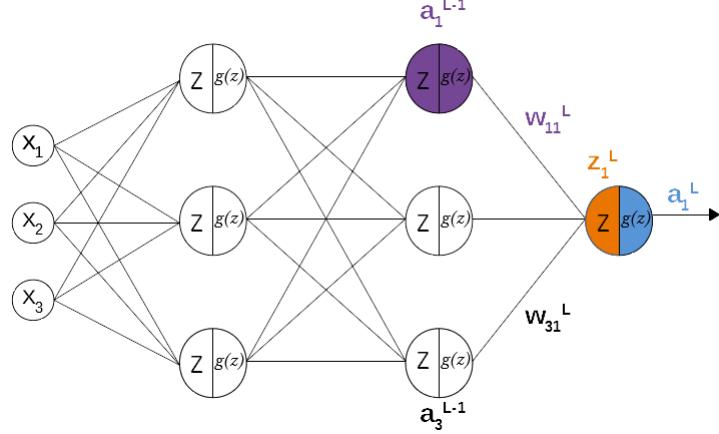
Figure 5: Neural Network with two hidden layers



The core principles are identical to the Perceptron, although one major modification has to be made. The gradient descent algorithm requires the ability to compute

all partial derivatives with respect to each weights. Because the first weights are passed through an activation function before reaching the next layers, we need this activation function to be differentiable⁵. In that case, we can apply the chain rule to compute partial derivatives from one end of the network to the other. Let us first define the notation.

Figure 6: simple network



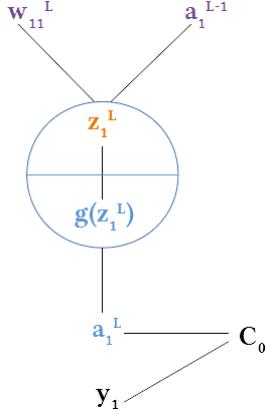
We associate each weight w to the two nodes it is connecting such that w_{jk}^L connects node j in the previous layer to node k in the current layer. The current layer is denoted as L and previous layers are denoted as $L-1, L-2, \dots, L-M$. We also define the loss function as C_n , $n = 1, 2, \dots, N$, N being the sample size and the activation function as g . The weighted sum of inputs arriving to node j is denoted as z_j^L and the output from node j is denoted as a_j^L . To use gradient descent we need to compute the gradient of the loss function to minimize. The gradient is therefore expressed as:

$$\nabla(C) = \begin{bmatrix} \frac{\partial C_n}{\partial w_{j,k}^L} \\ \dots \\ \frac{\partial C_n}{\partial w_{J,K}^{L-m}} \end{bmatrix} \quad (2)$$

We will show how to compute the gradient for one particular data point $n = 0$. We need to find the partial derivative of the loss function with respect to each weight in the network. For this example we will use the Mean Squared Error but other functions are appropriate to classification problems. This will be a two step procedure; first we compute partial derivatives with respect to weights on the last layer (output layer $L-1$). Then we generalize the results to every weight in the network. For simplicity, we assume the penultimate layer ($L-1$) contains only one neuron (a_1^{L-1}). The last two layers can then be represented as:

⁵We give further details on why that is when we compute partial derivatives

Figure 7: output layer



This diagram is inspired by the work of Grant Sanderson ([3Blue1Brown](#)) on [backpropagation](#).

The diagram shows how the output from node 1 in layer $L - 1$ is multiplied by a particular weight $w_{1,1}^L$ to obtain z_1^L , $z_1^L = \sum w_{i,j}^L * a_j^{L-1}$. z_1^L is passed through the activation function $g(z_1^L)$ and produces the output for node 1 in layer L , a_1^L . Finding partial derivatives with respect to weights on the last hidden layer $L - 1$ to the output layer L is straightforward because we can decompose the effect of weight $w_{1,1}^L$ with the chain rule:

$$\frac{\partial C_0}{\partial w_{1,1}^L} = \frac{\partial C_0}{\partial a_1^L} * \frac{\partial a_1^L}{\partial z_1^L} * \frac{\partial z_1^L}{\partial w_{1,1}^L} \quad (3)$$

The partial derivative of the loss function with respect to the weight $w_{1,1}^L$ is equal to the product of the partial derivative of the loss function with respect to the output of the last layer, the partial derivative of the output of the last layer with respect to the weighted sum z_1^L and the partial derivative of the weighted sum with respect to the weight $w_{1,1}^L$. Because the loss function is a function of the output given by the network, which is a function of the activation of the previous node which in turn, is a function of the weights, the partial derivative of the loss function is equal to the product of the partial derivatives of the functions that compose it. We yet have to compute these derivatives.

- $\frac{\partial C_0}{\partial a_1^L}$, the loss function is equal to: $(a_j^L - y_j)^2$. Therefore its partial derivative is: $2(a_j^L - y_j)$.
- $\frac{\partial a_1^L}{\partial z_1^L}$, the output a_1^L is the result of the activation function evaluated at z_1^L . We can also write; $a_1^L = g(z_1^L)$. Therefore its partial derivative is equal to the partial derivative of the activation function, $g'(z_1^L)$. This is the reason why we need differentiable activation functions such as the sigmoid.
- $\frac{\partial z_1^L}{\partial w_{1,1}^L}$, z_1^L is the weighted sum of the previous output multiplied by the weight $w_{1,1}^L$, a slight change in the weight will cause z_1^L to change by the value of the previous

output, a_1^{L-1} . The partial derivative of z_j^L with respect to the weight $w_{j,k}^L$ is equal to the previous output⁶.

We were able to compute one element of the gradient vector, for one particular training example:

$$\frac{\partial C_0}{\partial w_{1,1}^L} = 2(a_j^L - y_j) * g'(z_i^L) * a_1^{L-1} \quad (4)$$

Recall that we are using MSE to measure loss, MSE is the average error on all data points hence we need to compute a value for each data point $n = 0, \dots, N$. The gradient of the average loss is equal to the average gradient of the loss for each example therefore we repeat the process for each data point in the training set to obtain one element of the gradient vector which will be the average gradient for all points:

$$\frac{\partial \bar{C}}{\partial w_{j,k}^L} = \frac{1}{N} \sum_{n=1}^N \frac{\partial C_n}{\partial w_{1,1}^L} \quad (5)$$

This is only true for the partial derivatives with respect to weights on the last layer $L-1$, because the loss function is a direct function of the output from the last layer $C = C(a_j^L, y_j)$. However, it is not a direct function of the outputs from previous layers. This is where backpropagation is used, to rewrite $\frac{\partial C_0}{\partial a_j^{L-m}}$, $m \neq 0$ such that C_0 becomes a function of a_j^{L-m} using the chain rule. We want to solve:

$$\frac{\partial C_0}{\partial w_{j,k}^{L-1}} = \frac{\partial C_0}{\partial a_j^{L-1}} * \frac{\partial a_j^{L-1}}{\partial z_j^{L-1}} * \frac{\partial z_j^{L-1}}{\partial w_{j,k}^{L-1}} \quad (6)$$

$\frac{\partial a_j^{L-1}}{\partial z_j^{L-1}}$ and $\frac{\partial z_j^{L-1}}{\partial w_{j,k}^{L-1}}$ are derived exactly like if they were on the last layers as a_j^{L-1} , z_j^{L-1} and $w_{j,k}^{L-1}$ are direct functions of each other ($a_j^{L-1} = g(z_j^{L-1})$ and $z_j^{L-1} = \sum w_{j,k}^{L-1} * a_j^{L-2}$). We need to rewrite $\frac{\partial C_0}{\partial a_j^{L-1}}$ to be able to compute the gradient and identify the effect of weight $w_{j,k}^{L-1}$ on the loss function. Although the loss function is not a direct function of nodes that are not on the last layer, every node plays a role in the activation of all further neurons. Eventually, each neurons impacts the last output neuron, which in turn is a variable on which the loss function directly depends. The idea of backpropagation is to take advantage of the structure of the network to reverse engineer the propagation of the activation of neurons to be able to compute partial derivative of the loss functions with respect to all weights in the network. We would like to remind the reader that with all the partial derivatives we can calculate the gradient and optimize the weights in order to minimize the loss function. The first step to reverse engineer the network is to find $\frac{\partial C_0}{\partial a_j^{L-1}}$ by applying the chain rule on that expression:

⁶We use indices j and k as this is true for every weight

$$\frac{\partial C_0}{\partial a_j^{L-1}} = \sum \left[\frac{\partial C_0}{\partial a_j^L} * \frac{\partial a_j^L}{\partial z_j^L} * \frac{\partial z_j^L}{\partial a_j^{L-1}} \right] \quad (7)$$

We already demonstrated how to calculate $\frac{\partial C_0}{\partial a_j^L} = 2(a_j^L - y_i)$ and $\frac{\partial a_j^L}{\partial z_j^L} = g'(z)$. We will now focus on the new term:

$\frac{\partial z_j^L}{\partial a_j^{L-1}}$, the partial derivative of the weighted sum with respect to previous output is equal to the weight $w_{j,k}^L$ and $\frac{\partial C_0}{\partial a_j^{L-1}} = 2(a_j^L - y_i)g'(z)w_{j,k}^L$. Overall we were able to calculate the impact of the weights of the previous layer by applying a chain rule inside another chain rule:

$$\frac{\partial C_0}{\partial w_{j,k}^{L-1}} = \underbrace{\frac{\partial C_0}{\partial a_j^{L-1}}}_{2(a_j^L - y_i)g'(z)w_{j,k}^L} * \frac{\partial a_j^{L-1}}{\partial z_j^{L-1}} * \frac{\partial z_j^{L-1}}{\partial w_{j,k}^{L-1}} \quad (8)$$

This is where the idea of backpropagation stems from, to compute partial derivatives with respect to previous weights, one needs to compute the partial derivatives with respect to the last weights. The gradient must be calculated on the last layers first, and then the process can be expanded backward until the beginning of the network using the chain rule. The computation of partial derivatives is repeated for each data point and the average derivative is calculated $\frac{\partial C}{\partial w_{j,k}^L} = \frac{1}{N} \sum_{n=1}^N \frac{\partial C_n}{\partial w_{j,k}^L}$ before adding it to the gradient vector. After iterating through the entire data set, the gradient vector is used to update weights such that we take a step towards a local minimum of the loss function.

$$\overrightarrow{w^{new}} = \overrightarrow{w^{old}} + \eta(-\nabla(C))$$

The same stopping rules we discussed before are applied here to stop the gradient descent. We have seen how Neural Networks treat the information using a simplified example involving a Perceptron. We also explained how the model finds the optimal parameters by gradually minimizing the loss function by taking iterative steps following the gradient (inverse gradient). We then generalized the results to Neural Networks and showed how backpropagation can be used to solve the same problem with more complex networks. Until now we have considered vanilla examples, the following section deepens our understanding of various adjustments that can reinforce the network.

2.1.3 Activation functions and enhanced gradient descent

2.1.4 Support Vector Machine & the kernel trick

We are trying to find the maximum margin, in other words we want to maximize the width of the margin that is calculated as:

$$width = (X_+ - X_-) \frac{w}{\|w\|} \quad (9)$$

This is equivalent to maximizing $\frac{w}{\|w\|}$ or equivalently to minimize $\|w\|$. For mathematical we further modify the expression above to arrive at the following minimization problem: the largest margin is attained when $\frac{1}{2}\|w\|^2$ is minimum. This minimizing problem is constrained by: $y_i(X_i w + b) = 1$ as discussed earlier. Hence we are looking to solve:

$$\begin{aligned} & \min \frac{1}{2}\|w\|^2 \\ \text{s.t. } & y_i(X_i w + b) - 1 = 0 \end{aligned} \quad (10)$$

This is a standard constrained optimization problem. We will find the maximum margin using a Lagrangian defined as:

$$\mathcal{L} = \frac{1}{2}\|w\|^2 - \sum \alpha_i(y_i(w \cdot X_i + b) - 1) \quad (11)$$

To solve a Lagrangian we start by solving the partial derivatives with respect to each variable (α_i , b and w).

- $\frac{\partial \mathcal{L}}{\partial w} = w - \sum \alpha_i y_i X_i = 0 \rightarrow w = \sum \alpha_i y_i X_i$
- $\frac{\partial \mathcal{L}}{\partial b} = w - \sum \alpha_i y_i = 0 \rightarrow \sum \alpha_i y_i = 0$

We can plug these two equalities in \mathcal{L} :

$$\mathcal{L} = \underbrace{\frac{1}{2}(\sum \alpha_i y_i X_i)(\sum \alpha_j y_j X_j) - \sum \alpha_i y_i X_i \cdot (\sum \alpha_j y_j X_j)}_{= \frac{1}{2}a - a = -\frac{1}{2}a, \text{with } a = \sum \alpha_i y_i X_i} - \underbrace{\sum \alpha_i}_{=0}$$

We arrive at the final result:

$$\mathcal{L} = \sum \alpha_i - \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j X_i \cdot X_j \quad (12)$$

This result tells us that the optimization problem depends on the dot product of pairs of data points $X_i \cdot X_j$. To find the optimum, we know have to solve for α_i . Which means every pair of data points i and j will be weighted by some α_i . As we will discuss, this indicates that finding which $\alpha_i \neq 0$ is equivalent to finding support vectors. Once we find the support vectors (points on the marginal hyperplanes), we can easily find the intercept b and the slope of the decision boundary w .

At this point we would like to take a step back to illustrate the importance of the α_i term. Before we express w and b as sums of other variables we had:

$$\mathcal{L} = \frac{1}{2} \|w\|^2 - \sum \alpha_i \underbrace{f_i(w, b, X_i)}_{y_i(w \cdot X_i + b) - 1}$$

In order to maximize this expression with respect to α_i we would like $\sum \alpha_i f_i(w, b, X_i)$ to be as small as possible. We know from previous constraints that $\sum \alpha_i f_i(w, b, X_i) \geq 0$ hence, to maximize \mathcal{L} we need $\sum \alpha_i f_i(w, b, X_i) = 0$. This is possible in two different scenarios⁷:

1. $f_i(w, b, X_i) > 0$, which implies $\alpha_i = 0$
2. $f_i(w, b, X_i) = 0$, which implies $\alpha_i > 0$

These two scenarios tell us that for points X_i on the margin ($f_i(w, b, X_i) = 0$), α_i is strictly positive, for points not on the margin ($f_i(w, b, X_i) > 0$), $\alpha_i = 0$. Because w is defined as $\sum \alpha_i y_i X_i$, w only exists (w is not null) for points on the margin. For every other points, α_i is equal to 0 then, $\alpha_i y_i X_i = 0$ and w at that particular point i is null.

To maximize the margin, SVM only has to look at points on the margin, the other points will be irrelevant in the optimization problem. In other terms, once we solve for α_i , whenever $\alpha_i \neq 0$; the point is relevant and is on the margin. Solving for α_i is equivalent to finding points on the margin (support vectors), hence finding margins.

We then only need to differentiate \mathcal{L} with respect to α_i by setting the partial derivatives equal to 0:

$$\max_{\alpha_i} \mathcal{L} = \sum \alpha_i - \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j X_i \cdot X_j \quad (13)$$

We saw that optimizing this expression required that α_i is not null only for points on the margin. Therefore, once we set partial derivatives equal to 0, we will find optimal α_i which, plugged into $w = \sum \alpha_i y_i X_i$ will give us the w vector, orthogonal to the decision boundary with the largest margins.

We also know that $w^T X_i + b = y$, or $b = y - X_i w^T$. This strict equality is only true when X_i is on the margin. Hence for a subset s of i , $s \subset i$, for which X_s is on the margin, we have $b = y - X_s w^T$. Since we solved for X_s by solving for α_i we can compute b . The classification problem is solved, for any new point u we have:

$$\begin{aligned} &+ \text{ if } w \cdot u + b > 0 \\ &- \text{ if } w \cdot u + b < 0 \end{aligned} \quad (14)$$

w and b are known so we can classify every new point according on the hyperplane we found by solving for w and b .

⁷Both terms cannot be equal to 0

2.1.5 Soft Margin SVM

2.2 Finance

3 Application

We have defined all necessary notions in the previous sections, we will now apply the aforementioned techniques in order to predict asset prices. The following section is dedicated to machine learning and econometrics, we start with an exploratory analysis of the data and we proceed to model the information to hopefully capture the deterministic patterns in stock prices.

3.1 Data analysis

The data is obtained from the [Yahoo Finance API](#). We are using daily close prices adjusted for stock splits and dividends. We make the assumption that the market is not unique and different assets have different behaviors. Volatile assets are often considered as more reactive to technical indicators than very stable assets⁸. We can impute this empirical phenomenon to the fact that volatile assets attract speculative investors whom are prone to use technical indicators to time their trades. Whereas stable assets attract long term investors who might prefer large scale macro-economic variables to build their portfolio. We can already see how, considering this hypothesis, different variables are used for different markets. Technological assets, or stock from emerging markets which are notoriously volatile can be predicted with technical indicators whereas longstanding stocks such as Ford or Kellogg for instance do not rely on such indicators as much.

For this particular reason, we think it is reasonable to test our models on assets derived from different markets to asses whether or not these different behaviors can be highlighted. Nonetheless we will focus on US and Europe based companies, as exploring all possible assets is nearly impossible. We start by listing all the assets we are going to focus on in this study:

1. **Air France-KLM (AF.PA)** is a Franco-Dutch airline holding company born from the merge of two national airlines; Air France and KLM Royal Dutch Airline. Air France-KLM is one of the largest airline in the world and both French and Dutch governments are shareholders with approximately 14% each which makes it a very sound company. During the Covid-19 pandemic the company suffered from substantial losses due to the heavy decrease in air travel, funds have then been injected by the French government to help maintain Air France-KLM afloat.
2. **Tesla (TSLA)** is currently the most famous electric car manufacturer in the world. Tesla generates a lot of attention from technology enthusiasts and investors. The newly born company⁹ directly issued from the silicon valley has seen its market

⁸Stock return predictability: Evidence from moving averages of trading volume - Pacific-Basin Finance Journal 65 (2021) - Yao Ma , Baochen Yang, Yunpeng Su

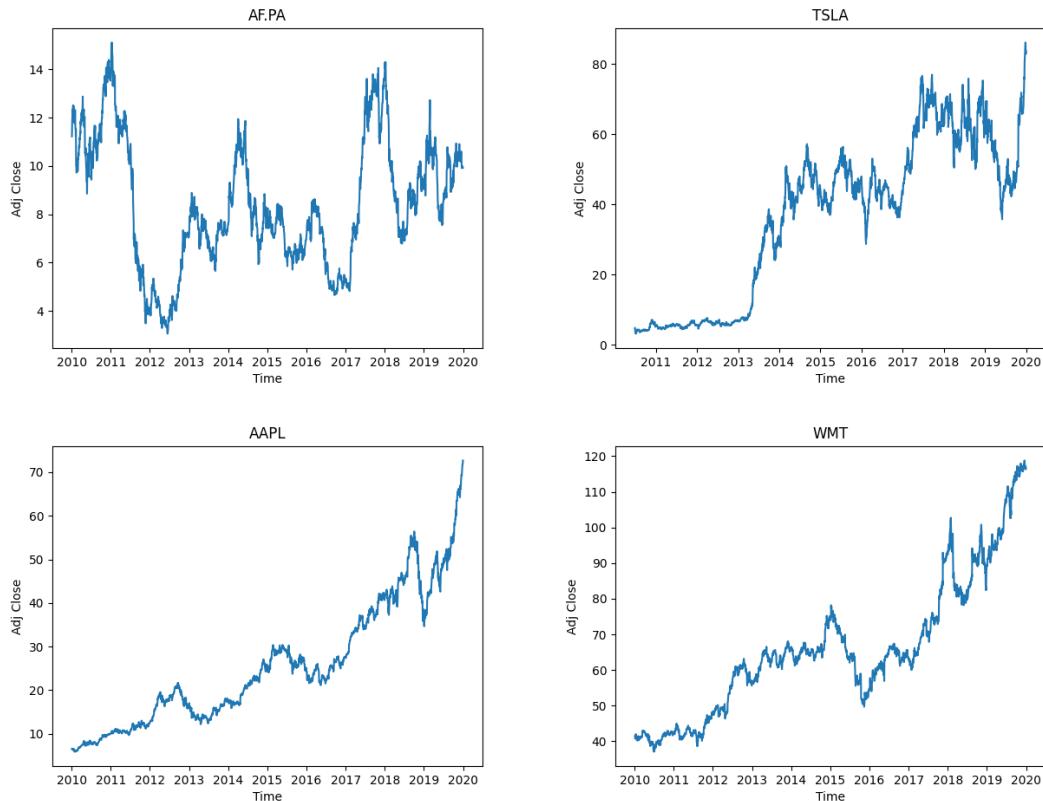
⁹Tesla was founded in 2003

value skyrocket since 2012. Some consider the valuation of the company does not reflect the true economic performance of Tesla, which is, along with speculative trading, one of the reason why Tesla's stock is famously volatile.

3. **Apple (AAPL)** combines both aspects, the company is innovative and has recorded the largest market capitalization in its history. But Apple is also considered as a sound investment because of its seniority and its deep implantation in the market for new technologies.
4. **Walmart (WMT)** is the largest retail corporation by revenue in 2020. Walmart is one the most notorious US-based firms in the world which makes it a good source of comparison with other assets.

We tried to choose these assets with respect to sectors, overall market sentiment on the valuation of the companies and volatility to reflect some of the disparities one can observe on financial markets. The data is gathered from 2010 January 1st to 2019 December 31st, for a total of approximately 2530 observations per asset. The period is chosen to reflect normal market conditions, post 2008 crisis and pre 2020 crisis. Although choosing this particular period help us get rid of abnormalities which makes the training more efficient, it does not reflect the entirety of the market conditions. Our model will therefore be incapable of predicting this kind of rare black swans events, which constitutes a major drawback to our approach.

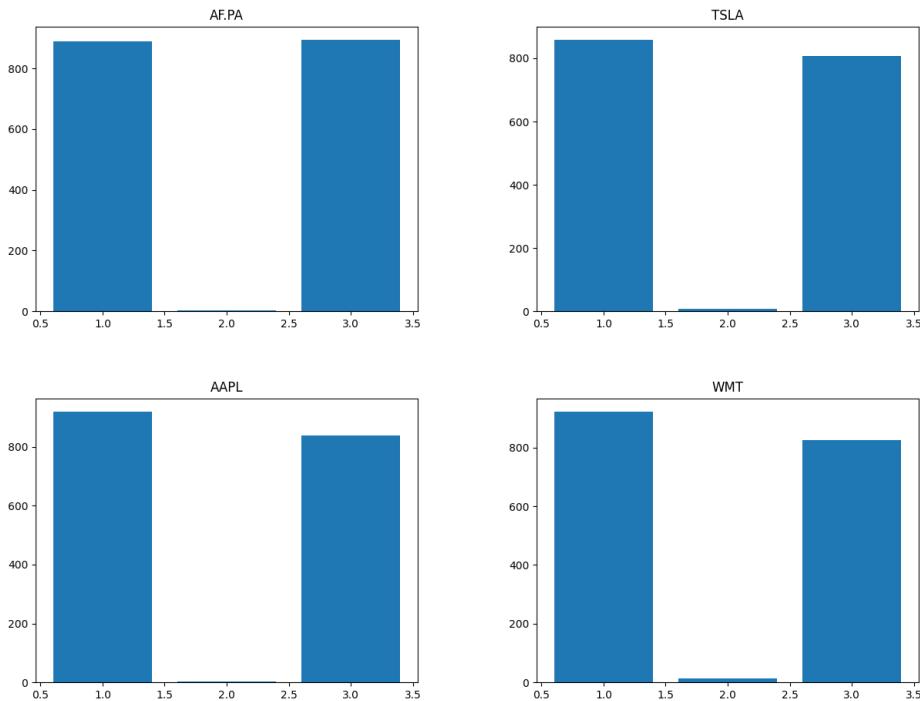
Figure 8: Series from 01/01/2010 to 31/12/2019



As we discussed in the previous section, we will be using classification models.

Therefore, we need to encode the close price into a categorical variable. We can create a new variable that takes a value of 1 when the next day price is superior to the current price, 2 if both prices are equal and 3 if the future price is below the current price. Although occurrences of the second case are very rare, we need to consider it. We will see later on if we can transform the problem such that we are able to use binary classifiers, which will be easier to interpret and more probably more efficient. The next figure represents the frequencies of appearance of each case for all four series.

Figure 9: Buy and Sell frequencies

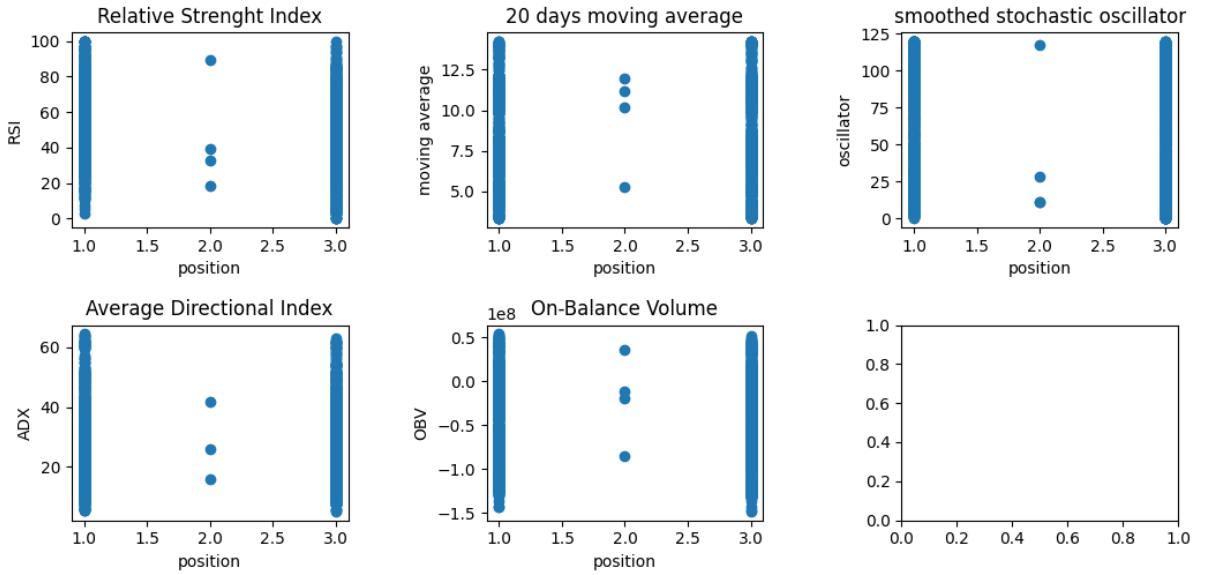


We can see how days where price is identical to the price the prior day are very rare. We fear that having this third category carries more noise than it carries useful information, that is the reason why we might consider transforming the problem into a 0/1, binary problem. Furthermore, buy and sell are represented equivalently in all four series with approximately 50% of the total number of days where returns were positive and the other 50% where returns were negative. This is one desirable property, as imbalanced categories induce a natural bias in terms of precision. If 90% of the sample belong to a certain class and 10% to the opposite class, even if the model assumes all observations belong to the first class, it will attain a 90% accuracy. Even with 90% accuracy, we can see how the model is not able to differentiate both classes. By only considering the percentage of correctly classified observations, we are encouraging the model to behave badly while giving a false impression of efficiency. Depending on how the model is trained, what loss functions are used and overall which error metric we choose, we can obtain very different results. In our case the buy (1) and sell (3) classes are very well balanced, which carries convenient properties, and further encourages us to treat days where the return is exactly 0.

In order to start our analysis, we visualize the **position** with respect to the

predictors in a two dimensional figure. This will help us determine whether or not the distribution of position is conditional on the predictors, which means the variable will change depending on the value of the predictor considered. In the other scenario, the **position** is independent from the predictor, meaning the value of the predictor on a certain day does not provide information on future price. Visual assertion of a relation between multiple variables is not sufficient to assume that variables are linked, we further need to assess the statistical significance of the relation.

Figure 10: Conditional distributions



Note: We only present a subset of the variables for one of the assets. The remaining figures are presented in the annexes.

At first glance, it seems the variables are uncorrelated with **position**. If we look at the RSI for example, we assumed that high values are associated with a 'sell' signal whereas small values are associated with 'buy' signals. However the RSI spans on the same range of values when we consider days where the price increased and days where the prices decreased, there is no distinct difference in the RSI depending on the right position. The same observations can be applied to the other predictors and can also be generalized to all four assets we are considering. This finding coincides with the EMH which states that since all the available information is already incorporated in the prices, technical indicators cannot anticipate future prices. Also, indicators are not supposed to be used on a daily basis, the theory behind it suggests that we should look for crossing points to detect trends, patterns or reversals as we explained in the last section. We will try to transform the variables such that instead of considering the absolute values we evaluate certain conditions to extract the true information from the indicators.

However we see no correlation between our variables vis a vis **position**, we need to statistically test for absence of significance of the variables to confidently establish that there is no residual information in the indicators. To fulfill this aim, we will need a model capable of generating hypothesis tests, which we cannot obtain with machine learning models. We will then try to classify **position** using a logit model. One important note

is that logit is a linear model. It is possible to have no linear dependencies between the indicators and **position** while still having non-linear dependencies between the variables. Although logit model will help us determine the nature of the relationships in our data set, it is limited by its linearity. We report the results in the table below

Table 1: Logit Regression

Asset : WMT		
LLR : $1.49e^{-16}$		
Variable	Coeff.	p-value
RSI	-0.0390	0.125

=====							
Dep. Variable:	position	No. Observations:	1742				
Model:	MNLogit	Df Residuals:	1726				
Method:	MLE	Df Model:	14				
Date:	Thu, 08 Apr 2021	Pseudo R-squ.:	0.1007				
Time:	18:18:36	Log-Likelihood:	-1144.4				
converged:	False	LL-Null:	-1272.5				
Covariance Type:	nonrobust	LLR p-value:	1.491e-46				
=====							
position=2	coef	std err	z	P> z	[0.025	0.975]	
RSI	-0.0390	0.025	-1.536	0.125	-0.089	0.011	
D	0.0230	0.013	1.726	0.084	-0.003	0.049	
MA	-1.718e+10	9.24e+05	-1.86e+04	0.000	-1.72e+10	-1.72e+10	
boll_up	8.591e+09	4.61e+05	1.86e+04	0.000	8.59e+09	8.59e+09	
boll_dw	8.591e+09	4.6e+05	1.87e+04	0.000	8.59e+09	8.59e+09	
MACD_short	-0.0463	0.015	-2.993	0.003	-0.077	-0.016	
MACD_long	0.0213	0.007	2.992	0.003	0.007	0.035	
adx	-0.0282	0.027	-1.036	0.300	-0.082	0.025	
OBV	-7.177e-09	3.69e-09	-1.948	0.051	-1.44e-08	4.59e-11	
position=3	coef	std err	z	P> z	[0.025	0.975]	
RSI	-0.0648	0.005	-13.526	0.000	-0.074	-0.055	
D	0.0250	0.002	10.174	0.000	0.020	0.030	
MA	-1.0882	2.19e+04	-4.97e-05	1.000	-4.29e+04	4.29e+04	
boll_up	0.5701	1.09e+04	5.21e-05	1.000	-2.14e+04	2.14e+04	
boll_dw	0.5459	1.09e+04	4.99e-05	1.000	-2.14e+04	2.14e+04	
MACD_short	0.0009	0.003	0.358	0.720	-0.004	0.006	
MACD_long	-0.0004	0.001	-0.360	0.719	-0.003	0.002	

Among all the assets, we were only able to estimate a logit regression for **WMT**. Other assets yielded NaN values for the coefficients and their standard error. We suppose this is due to very high multicollinearity among the predictors. Bollinger bands, which

are represented by **MA**, **boll_up** and **boll_dw**, represent the positive and negative limits of a two standard deviations interval centered around a moving average, they are constructed as a function of the moving average, hence, a completely dependent. Because of such multicollinearity, we cannot estimate the logit regression. Nonetheless, the estimation works with **WMT** data, but the MLE does not converge, and the multicollinearity is still present therefore the coefficients are biased and their standard error along with corresponding statistical tests are also biased. The first panel indicates the increasing or decreasing probability of switching from a buy position to a hold position. We can see how coefficients for **MA**, **boll_up** and **boll_dw** are very large with very small standard errors leading to a reject of the null hypothesis of no explanatory power¹⁰. These results further indicate presence of multicollinearity especially between these three variables. The smoothed stochastic oscillator **D** is significant at a 10% confidence level, both 20 days and 6 days moving averages **MACD_long** and **MACD_short** are significant at a 1% confidence level. Multicollinearity might also play a role in these two variables being significant. The On Balance Volume **OBV** is significant to a 10% confidence level and almost significant to a 5% confidence level. Relative Strength Index **RSI** is not significant neither is the Average Directional Index **ADX**. The Likelihood Ratio indicates the model is better fitter than the null model (p-value < 0.01), telling us our regression has at least some explanatory power on the position. The second panel shows the increasing or decreasing probability of switching from a buy position to a sell position. Here, **RSI** and **D** are the only two statistically significant variables. Both are significant to a 0.01% confidence level. The bollinger bands variables are not significant with very large standard errors. The moving averages are not significant. The different results further indicate the regression is biased, especially the fact that standard errors inverse when considering different positions. Despite the pitfalls in our model, we can believe that the stochastic oscillator and the RSI are still the two indicators that carry the most information on future prices as their values and values for their standard errors are the most rational.

As we saw earlier, technical indicators are often used as signals rather than as they are. Common practice is to establish thresholds which, once crossed upward or downward, send a buy or sell signal. Transforming the variables to replicate this behavior would first coincide with common practice and second would be more appropriate to our problem. As we saw, the raw values of the indicators are not dependent on the **position**, and carry very little if no information on future prices. Transforming the variables might restore some explanatory power. We will transform the variables and run another logistic regression to see if the estimation gives better results, and if the so transformed variable are better suited. In that case, we will try to correctly classify buy and sell opportunities using SVM and ANN and compare it to the logistic regression.

The transformations are consistent with existing practice, we briefly detail our procedure before moving on to the estimation.

- The RSI is divided by four different thresholds; crossing 70 upward, downward and crossing 30 upward and downward. We expect that crossing 70 upward send a sell signal and crossing 30 upward send a buy signal.
- The stochastic oscillator follows the same rules expect thresholds are replaced with 80 and 20 respectively.

¹⁰ $H_0: \beta_i = 0$ against $H_a: \beta_i \neq 0$, with β_i the coefficient of the i^{th} variable

- The bollinger bands are transformed such that crossing the upper bound sends a signal and crossing the lower bound sends another signal.
- The MACD are transformed such that when the short 12 days moving average crosses the long 20 days moving average upward the variable send a signal and when the short moving average crosses the long moving average downward sends another signal.
- OBV and ADX are kept as raw values.

We also transformed **position** into a binary variable to get rid of the second class with very limited representation across the dataset. We expect that removing this class will help the different classifiers as the data is not polluted by unnecessary information.

Figure 11: Logit regression with transformed variables

MNLogit Regression Results							
Dep. Variable:	position	No. Observations:	1666	Model:	MNLogit	Df Residuals:	1660
Method:	MLE	Df Model:	5	Date:	Sat, 10 Apr 2021	Pseudo R-squ.:	0.007704
Time:	15:44:13	Log-Likelihood:	-1145.9	converged:	True	LL-Null:	-1154.8
Covariance Type:	nonrobust	LLR p-value:	0.003219				
position=1	coef	std err	z	P> z	[0.025	0.975]	
RSI	-0.0938	0.043	-2.174	0.030	-0.178	-0.009	
D	-0.1345	0.058	-2.308	0.021	-0.249	-0.020	
MACD_long	0.3055	0.196	1.560	0.119	-0.078	0.689	
adx	0.0005	0.004	0.148	0.882	-0.007	0.008	
OBV	2.606e-09	1.04e-09	2.506	0.012	5.68e-10	4.64e-09	
boll	0.1627	0.137	1.189	0.235	-0.106	0.431	

MNLogit Regression Results							
Dep. Variable:	position	No. Observations:	1640	Model:	MNLogit	Df Residuals:	1634
Method:	MLE	Df Model:	5	Date:	Sat, 10 Apr 2021	Pseudo R-squ.:	0.008317
Time:	15:42:52	Log-Likelihood:	-1126.1	converged:	True	LL-Null:	-1135.6
Covariance Type:	nonrobust	LLR p-value:	0.002015				
position=1	coef	std err	z	P> z	[0.025	0.975]	
RSI	-0.0895	0.041	-2.162	0.031	-0.171	-0.008	
D	-0.0122	0.058	-0.210	0.834	-0.126	0.102	
MACD_long	-0.2704	0.184	-1.473	0.141	-0.630	0.089	
adx	0.0056	0.003	1.852	0.064	-0.000	0.012	
OBV	9.823e-12	7.2e-11	0.136	0.892	-1.31e-10	1.51e-10	
boll	0.4255	0.124	3.426	0.001	0.182	0.669	

MNLogit Regression Results							
Dep. Variable:	Position	No. Observations:	1640	Model:	MNLogit	Df Residuals:	1634
Method:	MLE	Df Model:	5	Date:	Sat, 10 Apr 2021	Pseudo R-squ.:	0.009676
Time:	15:46:10	Log-Likelihood:	-1123.7	converged:	True	LL-Null:	-1134.7
Covariance Type:	nonrobust	LLR p-value:	0.0005333				
position=1	coef	std err	z	P> z	[0.025	0.975]	
RSI	-0.1012	0.044	-2.287	0.022	-0.188	-0.014	
D	0.0062	0.056	0.111	0.911	-0.104	0.116	
MACD_long	-0.2709	0.188	-1.440	0.150	-0.640	0.098	
adx	0.0099	0.003	3.259	0.001	0.004	0.016	
OBV	2.864e-11	1.53e-11	1.870	0.062	-1.38e-12	5.87e-11	
boll	0.2515	0.133	1.895	0.058	-0.009	0.512	

MNLogit Regression Results							
Dep. Variable:	position	No. Observations:	1640	Model:	MNLogit	Df Residuals:	1634
Method:	MLE	Df Model:	5	Date:	Sat, 10 Apr 2021	Pseudo R-squ.:	0.006952
Time:	15:45:08	Log-Likelihood:	-1126.2	converged:	True	LL-Null:	-1134.1
Covariance Type:	nonrobust	LLR p-value:	0.007471				
position=1	coef	std err	z	P> z	[0.025	0.975]	
RSI	-0.0829	0.041	-2.001	0.045	-0.164	-0.002	
D	0.0591	0.053	1.109	0.267	-0.045	0.164	
MACD_long	-0.3365	0.195	-1.726	0.084	-0.719	0.046	
adx	0.0007	0.004	0.183	0.855	-0.007	0.008	
OBV	4.932e-11	3.36e-10	0.147	0.883	-6.09e-10	7.07e-10	
boll	0.4284	0.134	3.208	0.001	0.167	0.690	

From the previous table we notice a first improvement compared to previous attempt as the logit regression was estimated for all assets and all converged to the maximum likelihood. The likelihood ratio tests indicates all models have at least some explanatory power on their respective asset. However the pseudo R^2 indicate the models explain a very small fraction of the total variation of the **position** variable. Less than 1% for all models. From these first observations, we can assume technical indicators have very little correlation with future prices. The RSI is significant to a 5% level for all assets, the Bollinger bands are significant to a 1% level for **WMT** and **TSLA** and to a 10% level for **AAPL**. As mentioned in the earlier sections, we expected higher significance when considering "young" companies such as **AAPL** and **TSLA**, especially as both operate in the technological sector, usually more prone to speculative and short-term trading. The results seem to partially contradict our initial hypothesis. The ADX is significant for

AAPL and **TSLA** (to a 5% and 10% respectively) and not significant for **WMT** and **AF.PA**, which coincides with our initial statement. The stochastic oscillator is only significant for **AF.PA**, at a 5% level, and is not significant for every other remaining assets. The MACD is only significant to a 10% for **WMT** and is almost significant to a 10% level for **AF.PA**, **TSLA** and **AAPL**. The OBV is significant for **AF.PA** and **AAPL** (to a 5% and 10% levels respectively). Overall, results indicate that different assets react to different indicators, the Bollinger bands, the MACD and the RSI are the three indicators with highest significance overall. One reason to explain this would be that those particular indicators are among the most commonly considered by traders. It is responsible to assume that if a substantial fraction of investors on a market follow one indicator to place trades, the resulting supply and demand will cause the price to move accordingly. ADX seems to be significant when used on more volatile markets. The likelihood ratio also suggests similar observations, the fit of the model is best for **AAPL** and **TSLA** which also both have the two smallest likelihood ratio. We partially confirm that indicators play an higher role in relatively newer and unstable markets, but some observations contradict this hypothesis. Finally, even with very little explained variance, the models are all delivering useful information on future prices. With linear models, such as logistic regression, the predictions might not be salable to substantial economic gains, but maybe the inclusion on non-linearities is necessary to encompass the relationship between technical indicators and future prices.

We would also like to note that macroeconomic variables play a role in stock prices modeling, as well as many other variables such as weather, period of the year, business cycles etc... the omission of all determining variables is a pitfall in our model. First the coefficient of the selected variables are therefore biased (omitted variable bias), and second, our predictions will consequently be worse as we explain less of the variance. The inclusion of aforementioned types of variable definitely is a reasonable improvement to the model to consider. We can now train a Neural Network and a Support Vector Machine to classify "buy" and "sell" positions based on the transformed technical indicators. We will begin with a relatively network with two hidden layers of 10 neurons with relu activation functions and a final output layer of 2 neurons and a softmax activation function. Each of the final neurons will output the probability that the correct is "buy" first the first one and "sell" for the second, the final prediction then is the maximum of these two values. We use Sparse Categorical Cross Entropy as the loss function to optimize and the Adam optimizer¹¹. All parameters can and will be updated to improve the model. The performance will be assessed using the confusion matrix, precision, recall and f1-score. For clarity the results for each assets are presented separately.

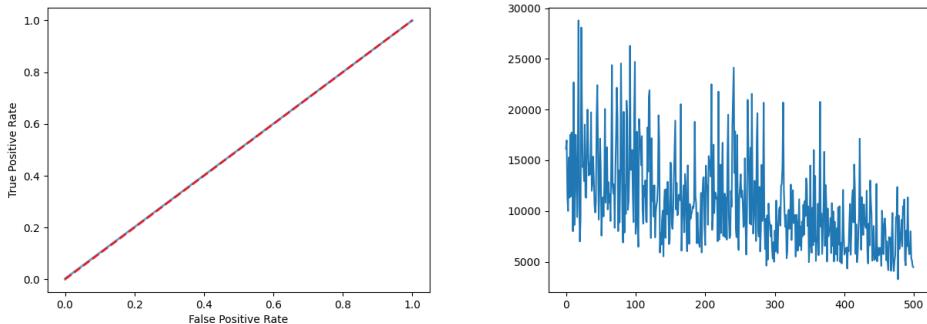
- **AF.PA**

Table 2: Confusion matrix

		True		Precision	0.50
		0	1		
Pred	0	0	0	Recall	0.25
	1	834	832		
		f1-score		0.33	

¹¹Adam is a stochastic gradient descent algorithm with adaptive learning rate which helps converge to global minimums more accurately and limits risk of overshooting

Figure 12: Results



The model associated 100% of the observations to the 1st class ("buy"). In other words, the Neural Net was not able to separate both classes and in results, associated every observations to one class almost at random. 832 out of 832 observations were correctly classified as "buy" which results in a 100% precision for that particular class, on the opposite 834 out of 834 observations that should have been classified as "sell" were also classified as "buy" which leads to a 0% precision for this second class. Overall, the total precision for both classes is 50%. Although 50% accuracy may seem satisfying, we see that in reality the model is not able to classify any observations correctly. The fact that 834 predictions are correct is due to the fact that the model associated all the observations to one class, each class represents roughly 50% of the data set, hence 50% of the sample are correctly classified. On the total number of observations (1666) associated to class 1, only 832 are relevant, resulting in a 50% recall for that class. On the total number of observations (0) associated to class 1, 0 are relevant, resulting in 0% recall for that class, the average recall on both classes then is 25%. The overall ability of our model to classify the data is given by the harmonic mean of precision and recall and is 33% in our case. The ROC curve illustrates the inability of the model to classify data. The curve lies on the median, telling us that the number of true positives is equal to the number of false negative, our model is wrong as many times as it is correct. The ROC curve also tells us that this is true for all possible thresholds. Here we did not specify a particular threshold and used 0.5¹², which is appropriate considering the balanced data set. The loss curve with respect to epochs shows a downward trend, suggesting the model is indeed learning from the data, and each iterations brings the loss closer to a minimum. However, the curve is very erratic and the downward trend is very mild which indicates the learning process is not smooth and constant as it should be with gradient descent. The impact of each predictor on the output is unclear and the gradient descent algorithm is therefore inefficient, even if it does indeed partially converges. We can notice the gradient descent is still converging after 500 epochs and has not reach a local minimum. To boost our model, one simple possibility would be to increase the number of epochs until the point where the loss stops decreasing. Using a binary loss functions and adding hidden layers to increase non linearity of the classifier may also be considered to improve the model.

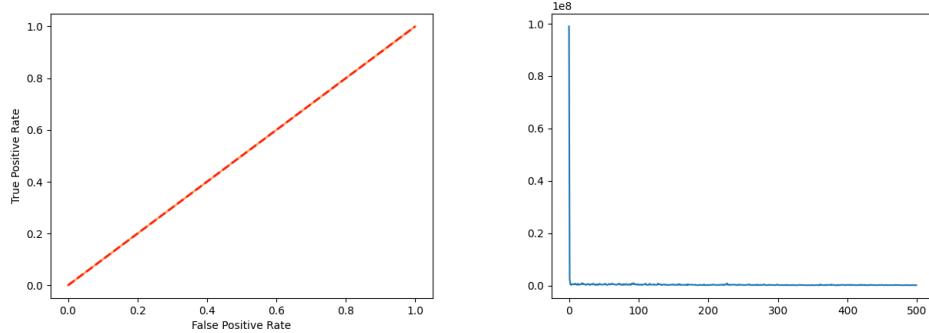
- **TSLA**

¹²If the probability that position belongs to class i exceeds 50%, it will be associated to class i

Table 3: Confusion matrix

		True		Precision	0.50
Pred	0	0	1		
	1	789	851		
				Recall	0.24
				f1-score	0.32

Figure 13: Results

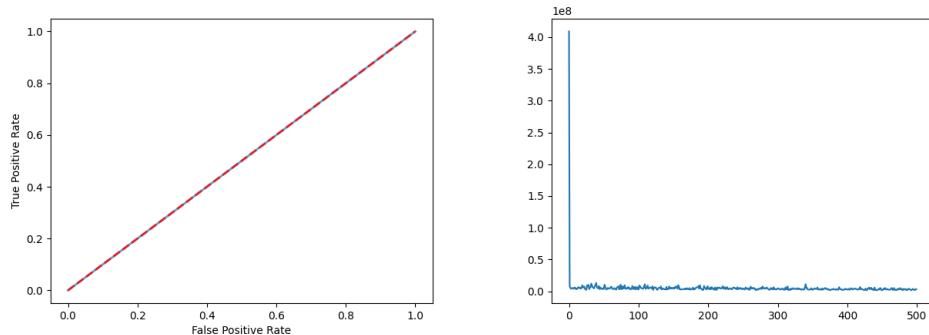


- AAPL

Table 4: Confusion matrix

		True		Precision	0.50
Pred	0	0	1		
	1	779	861		
				Recall	0.26
				f1-score	0.24

Figure 14: Results

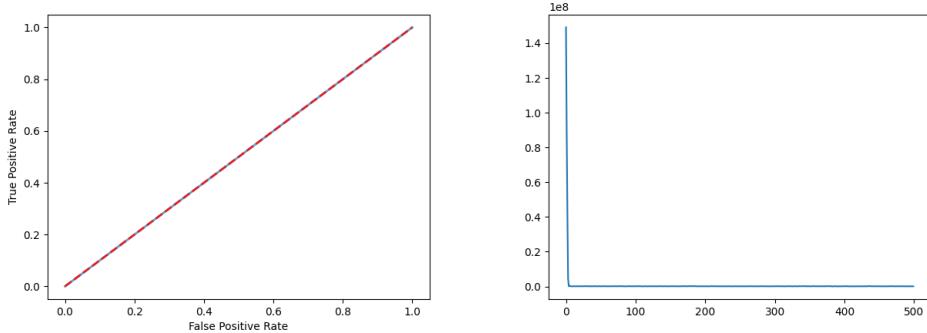


- WMT

Table 5: Confusion matrix

		True		Precision	0.50
Pred	0	0	1		
	1	773	867		
				Recall	0.24
				f1-score	0.32

Figure 15: Results



Similar observations are also true for the three other companies. The model is unable to classify the data, all observations are associated to one class at random. As the sample is constituted by roughly 50% of each class, we attain a 50% accuracy, that does not truly reflect the model's capacity to differentiate buy and sell opportunities. Those results are disappointing in a way but do not surprise us. As we saw earlier, the variables significantly help predict the future prices but we capture less than 1% of the total variance, we were not expecting very high performance. For **TSLA**, **AAPL** and **WMT** the gradient descent seems to converge very fast in contrast with **AFPA**. This result is surprising but could be attributed to chance, depending on the initialization of the different weights, the optimization can start more or less close to a local minimum. The convergences of the algorithm then depends on the initial weights vector. Convergence does not lead to good accuracy, we can imagine that considering the very small explanatory power of the variables, changing the weights does not induce substantial changes in accuracy. The gradients are very small and the loss functions is flat, the convergence to a local minimum then means convergence toward a point that is very slightly below the starting point. However the fact that the algorithm converges this fast (less than 10 epochs) seems almost unnatural given the results we first obtained with **AF.PA**. We did not find any particular errors that could cause this kind of behavior, we therefore assume this is due to the stochastic properties of Neural Networks. The ROC curves further indicate that even though the gradient descent converges, the final models are unable to classify the data.

We can now consider different approaches to optimize the model to improve classification accuracy. First, we increase the non-linearity by adding hidden layers and neurons. We will use 3 hidden layers of 32 neurons each. We can also increase the number of iterations to 1000 to ensure convergence and train the model for a longer time, on larger amounts of data. Less is often more, increasing the number of epochs might not be necessary, but it will be interesting to see the impact on the model's accuracy. Finally, we can choose a different loss function. We are currently using the Categorical Cross Entropy (CCE) which is a multi class probabilistic loss function. This function was appropriate before, as we had three different classes. However we transformed the data to avoid pollution due to very imbalanced data and we know face a binary classification problem. We then need a new loss function suited to our problem. We will then be using the Binary Cross Entropy (BCE) to try to improve our model's performance.

After testing different configurations, we were not able to improve the classifier. To this point, stock returns seem to be unpredictable which coincides with the Efficient Market Hypothesis. The model associates the entirety of the sample to either one of the

classes at random and even though the gradient descent appears to converge, accuracy metrics illustrates how poorly the model performs as in previous attempts.

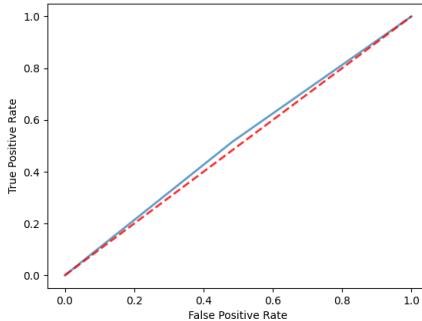
We will now train a Support Vector Machine on the same data to see if it performs better than the Neural Network. At first we will use a linear SVM with no kernel and proceed to add non-linearities with a gaussian kernel (Radial Basis Function kernel). The kernel hyperparameters will be optimized by grid search cross-validation.

- **AF.PA**

Table 6: Confusion matrix

		True		Precision	0.52
		0	1	Recall	0.52
Pred	0	431	403	f1-score	0.52
	1	403	429		

Figure 16: ROC



At first glance the results seem to indicate a better training, the area under the ROC curve is positive, the ROC goes slightly above the median. The ratio of True Positives (TP) to False Positive (FP) is positive, meaning the model is more often correct than it is incorrect. The confusion matrix indicates $431 + 429 = 860$ observations where correctly classified against $403 + 403 = 806$ misclassified observations. The precision and ratios are therefore greater than 0.5, because for both classes the number of correctly classified observations is superior to the number of misclassified ones (causing precision to increase), but also because among all the observations considered as belonging to one particular class by the model, more observations were really belonging to that class compared to observations that belong to the other class (causing recall to increase). Both precision and recall equal 0.52 which leads to a similar f1-score.

Although those results appear to be positive, the linear SVM does not perform better than the Neural Network. The difference in accuracy, is very small, probably negligible. Even more importantly, the SVM adopted the same strategy as the NN. Being unable to classify the data, it associated 50% of the data to one class and 50% to the other class, where NN associated 100% of the sample to one class. We see how both strategies lead to the same result, with perfectly balanced data, the SVM

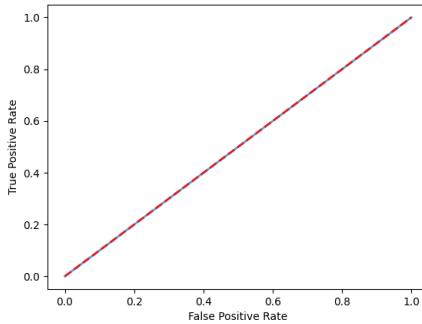
would have gotten similar identical results; 50% precision, 50% recall¹³ and a null area under the ROC curve. Here, because the data is not perfectly balanced, when obtained 52% overall accuracy, but the model does not outperforms the NN, both models are incapable of classifying the data. We will train the same model on all assets, if we observe similar results, we can conclude standard Neural Networks and linear SVMs trained on technical indicators seem to prove the EMH. We can see how 50% accuracy of the model connects to the hypothesis under which stock prices follow a random walk. If at any given time we cannot know if a certain stock price is going up or down and place a bet on either one direction, it is similar to a coin toss, where the probability of each side is 50%. If we repeat the experiment a sufficient number of times, we will be right 50% of the time and wrong 50% of the time, as the law of large number suggests.

- **TSLA**

Table 7: Confusion matrix

		True		Precision	0.50
Pred	0	0	0	Recall	0.26
	1	789	851	f1-score	0.34

Figure 17: ROC



4 Conclusion

¹³ Associating 50% of the data set to each class leads to 50% recall for both class. With Neural Network associating 100% to one class leads to 50% recall in one class and 0% in the other the overall recall then is 25%. SVM chose a "less risky strategy", but results are similar.